

EnterprateAI Funding Readiness Product Requirements

Version 0.1 | 30 September 2026 | Product and engineering review

1 Purpose and product outcome

Funding Readiness Check helps business owners understand how prepared they are to approach a selected funding route, identify evidence and preparation gaps, and complete the next useful actions. It reuses EnterprateAI business information, intelligence engines, scenarios and document capabilities.

The feature must produce an explainable readiness assessment, a funding requirement and use-of-funds record, an evidence checklist, prioritised actions and an investor preparation pack. Readiness measures preparation against the published criteria. It does not represent the likelihood of securing investment or replace an investor's assessment.

This PRD implements the Shared Architecture and Feature Integration Blueprint v0.1. Shared permissions, assessment history, data provenance, event handling and agent controls remain governed by that blueprint. Production mapping is an engineering prerequisite: documented entities and services must be checked against deployed code before migrations are specified.

2 Users and release scope

User pathway	Need	Required behaviour
Startup	Prepare a credible funding case before or during early revenue	Accept forecasts and research evidence with clear labels; do not require invented trading history
Small Business	Prepare an expansion or other equity funding case	Reuse operating records and distinguish historical performance from projections
Multi-Entity	Review preparation across authorised businesses	Keep cases and scores entity-specific; permit drill-down without averaging readiness scores

Proposed first release: UK service and software businesses preparing for equity discussions. This is a recommended launch boundary requiring product confirmation, not an established user restriction. Store route, jurisdiction, business model and stage as independent configuration dimensions. Debt and grant assessment profiles follow after their criteria and evidence requirements are separately defined. Unsupported profiles show availability status and cannot produce a misleading equity-derived score.

MVP includes case creation, prefill and confirmation, manual evidence entry and uploads, a 12-month monthly cash forecast input or mapping, deterministic assessment, action tracking, assessment history, dashboard summary and an exportable readiness report. Investor material preparation includes an editable executive-summary draft and structured pitch-deck content, reusing existing generation capabilities where available.

Later releases include lender-specific criteria, grant programme profiles, advanced document extraction, automated evidence refresh, full presentation design, investor matching and supervised funding agents. MVP does not send applications, contact investors, value the business, recommend securities or determine SEIS/EIS eligibility. Any tax-status document is stored as user-provided evidence without inferring eligibility. Endorsement assessment remains a distinct future profile.

3 User journey and screens

1. Enter Funding Readiness from navigation, a dashboard card or a Business Plan Generator action. Select the business first; persist that scope throughout the journey.
2. Create or resume a funding case. Enter route, target amount, currency, intended receipt date, purpose and intended milestone. Allow an unknown amount initially and offer the forecast worksheet to help investigate it.
3. Review prefilled profile, offerings, financial inputs, demand evidence and relevant engine outputs. Show source and date beside each reused item. Ask only for missing or conflicting information.
4. Complete six guided sections: funding goal; business and market; traction; finances and use of funds; team and delivery; documents and risks. Save drafts automatically after valid edits and provide explicit retry on save failure.
5. Run the check. Show progress and allow navigation away while the job continues. Present a preparation classification, evidence coverage, criterion results, blockers and the three highest-priority actions.
6. Open an action to view why it matters, the evidence required and the relevant existing tool. Update the underlying record or attach evidence, then reassess.
7. Prepare and export materials from a selected assessment snapshot. Show incomplete claims as gaps and retain a traceable link to source facts.
8. Return to the case to track change. Explain whether a changed result came from new evidence, changed assumptions or a new ruleset.

Screens: case list; guided setup; evidence and financial review; assessment results; action detail; materials; history. On mobile, use sequential sections rather than a wide spreadsheet layout. The result screen must support keyboard operation, readable status text and non-colour status cues.

4 Functional requirements

ID	Requirement	Priority
FR01	Create multiple named cases per business; archive without deleting historical assessments	MVP
FR02	Prefill authorised data with provenance; preserve explicit case assumptions separately	MVP
FR03	Record funding target, timing, use of funds and measurable milestones	MVP

ID	Requirement	Priority
FR04	Attach evidence to claims and criteria with source, effective date and verification status	MVP
FR05	Capture a monthly cash forecast and distinguish actual, forecast and estimated inputs	MVP
FR06	Calculate versioned readiness with blockers, missing data and traceable criterion results	MVP
FR07	Create deduplicated actions with owner, status, due date and completion evidence	MVP
FR08	Generate a readiness report, executive-summary draft and pitch-content outline	MVP
FR09	Expose scoped dashboard summaries and immutable assessment history	MVP
FR10	Run a funding timing or amount scenario using existing simulation infrastructure	MVP subject to validated engine support
FR11	Delegate evidence collection and draft preparation through registered agent tools	Later

FR10 is an explicit release decision: if the simulation service is not ready, release the assessment with a clearly unavailable scenario action and record FR10 as deferred. Do not replace simulation with an LLM prediction.

Case lifecycle: draft, active, archived. Assessment execution lifecycle: queued, running, succeeded, failed, cancelled. Readiness classification and execution status are separate. Archived cases are readable and exportable by authorised users but require reactivation before edits or a new assessment.

5 Data model and ontology integration

Structure	Reuse or addition	Required information
Business and BusinessProfile	Reuse Identity domain	Business ID, country, stage, model, currency and profile version
FundingCase	Add within Decision and Simulation domain	ID, business ID, route/profile version, title, target amount, currency, timing, purpose, owner, status, revision
UseOfFundsItem	Add linked case detail	Category, amount, period, milestone reference and forecast cash-line reference
FundingMilestone	Reuse shared milestone structure or add scoped record	Outcome, measure, baseline, target, due date, owner and linked allocation

Structure	Reuse or addition	Required information
EvidenceItem and EvidenceLink	Reuse shared evidence service	Claim/criterion reference, source record or attachment, effective date, verification and expiry
Financial records and forecast	Reuse Financial domain and plan forecast service	Monthly receipts, payments, opening cash, assumptions, basis and version
Assessment and result entities	Extend existing assessment service	Case revision, input snapshot, engines/rules versions, criterion results, flags and recommendations
DocumentArtifact	Reuse document service	Type, artifact ID, case/assessment snapshot, version, generated date and review status

A funding case owns its requested amount, allocations and case assumptions. Canonical profile and transaction records remain owned by existing services. Editing a case assumption does not update live business records. An explicit update action must describe the affected canonical record and require the appropriate write permission.

If a reusable forecast structure is absent, add a versioned Forecast and ForecastLine within the Financial domain. ForecastLine identifies month, category, cash direction, amount, currency, assumption and source. It must support baseline and scenario references and must never write forecast amounts into actual revenue or transaction tables.

Every new record is business-scoped. Foreign-key or equivalent service validation must prevent references to another business's evidence, documents or assessment. Use decimal money values or integer minor units, ISO currency codes and explicit period boundaries. MVP calculations operate in one case currency; mixed currencies require explicit conversion provenance before assessment.

6 Financial input and reconciliation rules

The proposed MVP horizon is 12 consecutive monthly periods starting from an explicit forecast date, extended when a selected milestone is later. This is a product preparation requirement, not a universal investor requirement. Pre-revenue users may enter zero actual revenue; blank means unknown, not zero.

Monthly closing cash equals opening cash plus cash receipts plus financing receipts minus cash payments. The next month's opening cash equals the prior closing cash. Store financing separately from trading receipts. Use-of-funds allocations reference forecast payment lines so the same expenditure is not counted twice.

Proposed funding is included only in the funded scenario at its assumed receipt date. The baseline excludes the proposed raise. Already available cash belongs in opening cash; committed but unreceived funding is separately dated and evidenced. Soft indications of investor interest are not cash or committed financing.

Reconcile target funding to allocations, including explicit reserve and fundraising-cost categories where relevant. In the case currency, allocations must equal the requested amount to the smallest supported currency unit before the allocation criterion is met. Do not silently change the target to make the totals match.

Show baseline minimum cash, first negative-cash month, funded minimum cash, funding timing gap and milestone coverage. A cash gap before the assumed funding date remains visible even if later funded balances are positive. Monthly resolution cannot prove that every intra-month payment can be made; label the projection as monthly.

Reuse validated engine computations for contribution margin, capacity, cash burn and concentration. Do not reuse the source viability document's break-even-months label until corrected. Resolve the blueprint's formula/example discrepancies before any affected metric becomes a readiness input. Unknown or unbounded runway is a typed state, never a JSON infinity or an invented number.

7 Proposed assessment criteria

The following pilot rubric measures preparation. Weights and thresholds are product proposals requiring calibration before public release. None is an empirically validated predictor of funding outcomes. Engine scores inform explanations and relevant checks; they are not simply averaged into readiness.

Each criterion below has two checks. Store the result of each check and its evidence. Score the criterion as 100 when both checks pass, 50 when one passes and one fails, and 0 when both fail. Any unknown required check makes the criterion unknown. A failed gate overrides any positive aggregate result.

Criterion and weight	Check A	Check B
C1 Funding purpose 10	Positive target, currency and intended receipt date confirmed	Purpose links to a measurable milestone
C2 Use of funds 15	Allocations reconcile exactly to target	Every allocation has a period and purpose or reserve rationale
C3 Financial preparation 20	Required monthly forecast is complete and reconciles	Material receipts and payments have documented assumptions and sources
C4 Demand evidence 15	At least one dated customer or market evidence item is attached	Its relevance to the target customer and offer is confirmed by a reviewer
C5 Business economics 10	Pricing and direct/fixed costs are provided on a consistent basis	Break-even/capacity findings and limitations are reviewed and explained
C6 Market and growth plan 10	Target segment and differentiation are documented	Acquisition approach has owner, cost assumptions and measurable milestone
C7 Team and delivery 10	Key execution responsibilities have named owners	Material capacity and dependency gaps have documented actions
C8 Materials and risks 10	Required preparation documents exist and are reviewed	Known material risks and funding-timing issues are disclosed with actions

Required MVP preparation documents: executive summary or equivalent pitch content, use-of-funds schedule, monthly forecast and team/ownership summary. Existing documents can satisfy requirements; generation is optional. Ownership information describes preparation and is not a legal verification. Qualitative checks require

explicit reviewer attestation and a supporting rationale; an LLM suggestion cannot independently mark them as passed.

Trading businesses provide demand evidence appropriate to their claimed traction. Pre-revenue businesses may use research, enquiries or pilots; the evidence type and strength must remain visible. Document completion does not prove commercial attractiveness. For this first profile, all eight criteria apply; unsupported business models must not bypass them through user-selected not-applicable flags.

8 Scoring classification and confidence

When all criteria are known, readiness score is the sum of each criterion score multiplied by its percentage weight. The weights total 100. Round only the displayed score to the nearest whole number; evaluate thresholds on the unrounded value. When any criterion is unknown, return a null overall score and show weighted evidence coverage plus completed criterion results.

Apply classification in this order:

1. A confirmed failed gate gives `not_ready`, even if other evidence is unknown.
2. Otherwise any unknown required check gives `insufficient_evidence`.
3. Otherwise a score below 60 gives `not_ready`.
4. Otherwise a score from 60 to below 80 gives `conditionally_ready`.
5. Otherwise a score of 80 or above gives `ready`, provided all gates pass.

Display `ready` as “Prepared for initial equity discussions against this checklist.” Display `conditionally_ready` as “Preparation gaps remain.” Show that business risk, evidence strength and readiness are separate dimensions.

Gates: G1 a confirmed positive funding target and explicit timing; G2 reconciled allocations; G3 complete and internally reconciled forecast; G4 no unresolved contradictory material figures across the selected case documents. A missing gate input is unknown. A confirmed invalid or conflicting input fails the gate. For G4, deterministic checks cover target amount, currency and forecast period; other contradictions can be raised by a reviewer. Unknown documents leave G4 unknown until reviewed.

Evidence coverage is the percentage weight of criteria for which both checks are known. Confidence is reported separately: low if any material source is missing, conflicted or stale; medium if complete and current but material claims rely on self-attestation; high only if complete/current and each material historical claim is linked to corroborating records or a documented independent review. Forecasts remain assumptions even when well evidenced. A connected system alone does not confer high confidence.

Pilot freshness defaults: cash position 30 days; operating results through a period ending within 90 days; current-case forecast confirmed within 30 days; team, market and demand evidence reviewed within 180 days. Explicit expiry dates take precedence. Record these policies in the profile version. Stale material evidence becomes unknown until refreshed or a qualified reviewer records a reasoned exception. Reconfirmation must not falsify the original evidence date.

9 Results actions and investor materials

The result includes case name, route, assessment date, data basis, classification, score when available, coverage, confidence, dimension results, gate status, business risks and next actions. Every finding links to a criterion, input and ruleset. Historical results remain readable after later data changes.

Rank actions deterministically: failed gates first, missing mandatory evidence next, then other gaps. Within each group use deadline, descending criterion weight and stable action ID. Explain the priority and avoid promising a score increase before reassessment. Repeated runs update or link an existing unresolved action for the same case and criterion instead of generating duplicates.

Action states: open, in_progress, awaiting_evidence, completed, dismissed. Completion requires evidence or a recorded reviewer decision. Dismissal records a reason and does not alter the underlying criterion.

Reassessment, rather than checking a box, updates readiness.

Readiness report export contains methodology version, findings, limitations, sources and actions. Executive-summary and pitch-content drafts draw facts from the selected snapshot and visibly label forecasts. Missing claims appear as named gaps. Never invent clients, traction, approvals, market statistics or commitments. Generated artifacts use draft/reviewed/superseded status, and retain artifact-to-source references.

A changed funding target or forecast marks dependent drafts as needing review. An archived export remains an as-of-date artifact and must not be silently rewritten. In-app downloads require current access checks. Public sharing and automated investor delivery are outside MVP.

10 Dashboard and agent interfaces

Startup dashboard: funding preparation card shows the selected case, classification, evidence gaps and next action. Small Business dashboard: the same card adds relevant cash and operational risks. A user may pin a case; otherwise select the most recently updated active case and make the selection clear.

Multi-Entity dashboard lists authorised businesses with case name, route, readiness status, freshness and critical gaps. It never computes an average readiness score. Compare numerical scores only within the same profile/ruleset, and display differences in stage and evidence basis. Pending or failed runs do not appear as zero scores.

Publish a stable summary contract for the Adaptive Dashboard PRD: `business_id`, `case_id`, `assessment_id`, `profile_version`, `classification`, `score`, `coverage`, `confidence`, `freshness`, `blocker_count` and `next_action_id`. Unknown values remain null. The assessment service owns calculation; the dashboard only presents it.

Funding-related future tools include `get_funding_case`, `list_readiness_gaps`, `propose_case_update`, `run_funding_assessment`, `prepare_funding_summary` and `simulate_funding_change`. Read tools, draft writes and external actions require different permissions. MVP works without an agent. An agent may prepare updates but cannot silently overwrite canonical records, attest evidence as verified or submit a funding application.

11 Permissions APIs and reliability

Map capabilities to existing roles rather than inventing a second role system. Required capabilities: `funding.read`, `funding.edit`, `funding.assess`, `funding.review`, `funding.export` and `business.membership.manage`. Viewing a business does not automatically confer access to every confidential attachment. Reviewer attestation records actor, time, evidence and rationale.

Proposed operations: create/list/update/archive funding cases; add allocations and evidence links; request/read an assessment; list/update actions; generate/read/download materials. Route them through existing authenticated business-scoped APIs. Updates require a revision token; stale writes return a conflict with a recoverable comparison. Assessment requests include case revision and idempotency key. Background jobs return a run ID and expose status.

Persist immutable input snapshots and selected evidence versions before running calculations. When input changes during a run, retain the result as historical and mark it stale; never present it as the current assessment. Publish `FundingCaseChanged`, `EvidenceUpdated`, `AssessmentCompleted` and `AssessmentInvalidated` through the shared reliable event mechanism. Coalesce recalculation requests and make consumers safe under duplicate delivery.

Protect uploaded files with existing size/type checks, malware handling, access controls and retention policies. Reject cross-business references at the service boundary. Check access again on background execution and download. Avoid raw financial values, documents or personal information in analytics and operational logs.

Proposed performance budgets for engineering validation: 95th-percentile case reads below two seconds; assessment results within ten seconds after inputs are available, excluding narrative generation; document jobs asynchronous with progress and timeout. Set load assumptions before accepting these targets. LLM failure must leave a deterministic report usable. Failed document jobs can retry without creating duplicate user charges or contradictory versions.

12 Acceptance examples

ID	Given and when	Expected result
AT01	All eight criteria score 100 and gates pass	Score 100; ready; forecast labels remain visible
AT02	All criteria score 50; gates pass	Score 50; not_ready
AT03	C3 scores 50; all others 100; gates pass	Score 90; ready with C3 gap disclosed
AT04	Target is 100000; allocations total 90000	G2 fails; not_ready regardless of score
AT05	C4 unknown; all others known and no gate fails	Null score; 85 percent coverage; insufficient_evidence
AT06	C4 unknown and G2 fails	not_ready; missing evidence also shown
AT07	Pre-revenue user enters zero sales and documented research	No fabricated customers/history; evaluate submitted evidence and forecasts
AT08	Proposed funds arrive after first negative-cash month	Show timing gap; do not treat proposal as current cash
AT09	Case changes while assessment runs	Result retained with original snapshot and stale status
AT10	Evidence expires after a successful assessment	Current view becomes stale; reassessment handles affected checks as unknown

ID	Given and when	Expected result
AT11	Another entity ID is supplied in URL, attachment or job	Access denied; no details leak into summaries or exports
AT12	Assessment or export request is retried	One logical run/artifact per idempotency key; no duplicate charge
AT13	LLM unavailable	Deterministic criteria, classification and actions still returned
AT14	Action is dismissed without resolving evidence gap	Criterion remains unchanged
AT15	Funding simulation executes	Live balances, expenses and assessments are not overwritten
AT16	New profile version is released	Existing history remains reproducible; cross-version comparison is labelled

Add boundary tests at exactly 60 and 80, blank versus zero, negative amounts, duplicate evidence, invalid dates, unsupported currencies, revoked access and conflicting document facts. AT03 requires the failed financial check to be assumption documentation, not the forecast-reconciliation gate.

13 Delivery plan and release decisions

Build in five work packages: production mapping and rubric review; case/evidence/forecast persistence; assessment and action service; guided UI/dashboard/export; pilot calibration and staged rollout. Establish entity isolation, versioning and input validation before user-facing generation.

Use additive migrations and feature flags. Backfill only documented source mappings and retain provenance. Do not translate an old score into the new rubric without the necessary evidence. Validate existing quotations, invoices, contracts and dashboard flows during regression checks. Disable new entry points for rollback while retaining user drafts and assessment history.

Pilot with both pre-revenue and trading businesses. Have business reviewers independently assess the same cases, examine disagreement and identify criteria users cannot answer. Revise the versioned rubric before public scoring. Measure usefulness and resolved gaps; do not optimise for artificially high scores.

Track: assessment completion rate; time from case creation to first actionable result; percentage of required fields reused; evidence coverage; proportion of cases resolving a blocker within 30 days; repeat assessments; draft correction rate; assessment failures; cost per successful run. Baseline these during pilot before committing commercial targets.

Release decisions and owners: Product confirms initial route/sector scope, subscriptions and usage limits; Product with domain reviewers approves rubric and public wording; Engineering confirms forecast support, engine corrections, role mapping and performance budgets; Product and Engineering confirm simulation inclusion; Security owner validates entity/document access and retention behaviour. These decisions gate the affected release, not drafting or prototyping.

14 Architecture traceability and handoff

Reuse the blueprint's AD-01 through AD-09. In particular, keep readiness as an assessment/service layer within the existing ontology, preserve deterministic ownership, maintain immutable snapshots and separate agent authority from recommendations. FundingCase is the subject of a readiness assessment; Business remains its owner.

Reference specifications: EnterprateAI Shared Architecture and Feature Integration Blueprint v0.1, sections 4 to 6 and 9 to 13; Updated Business Ontology, assessment/evidence/financial/scenario entities; Survival, Stability, Service Idea Viability, Growth and Real-Time Simulation engine specifications. The blueprint's source-validation findings are prerequisites to reuse of affected calculations.

The UK Endorsing Services materials are not funding-approval rules for this profile. External programme rules, investor-specific criteria and tax eligibility require separately maintained and verified profiles before they can affect a production assessment.

Engineering handoff must provide the implemented-to-proposed field mapping, API contracts, migration plan, versioned rubric configuration, representative fixtures for every classification, permission matrix and rollout checklist. Design handoff must include guided sections, empty and error states, missing-evidence results, business switching, document review and an accessible mobile layout.