

ENTERPRATEAI

Goal-Driven Conversational Onboarding & First-Value Activation

Product Requirements Document

Version 2.0 - Draft for Product & Engineering Review

Adopted product principle Do not onboard the user into software. Let the user start a task, and let EnterprateAI build the onboarding around that task.

7 October 2026

1. Document Control

	Value
Document ID	PRD-GO-001
Project Name	Goal-Driven Conversational Onboarding & First-Value Activation
Version	2.0 - Draft for Review
Status	Proposed Version 2 build baseline; Version 1.0 remained unadopted
Product	EnterprateAI
Release Date	7 October 2026
Parent Architecture	EnterprateAI Shared Architecture and Feature Integration Blueprint v0.1
Related PRDs	Adaptive Dashboard v1.0; Agentic Orchestration v1.0; Marketplace Business Index & Claim GTM v0.1; Launch Readiness v0.1; Funding Readiness v0.1
Product Owner	Product / Enterprate Limited
Primary Focus	Outcome-led homepage discovery: business problems/goals are primary entry points, quick tasks remain secondary, and all routes converge through the same Goal Resolver and workspace.

2. Executive Summary

EnterprateAI will replace traditional setup-first homepage onboarding with an outcome-led, goal-first conversational experience. The homepage intentionally provides two visible entry modes: prominent business goals/problems such as “I need funding”, “I need more customers”, “Help me price my product/service”, “I want to launch a product/service”, “Improve my cash flow”, “Help me grow my business”, “Test a business idea” and “Help me understand my business risks”; and a secondary quick-task path such as Create Invoice, Create Quotation, Create Business Plan and Prepare Proposal. A visitor may also type any goal in natural language. EnterprateAI resolves the underlying business outcome, validates it against registered product capabilities, and then starts a direct task, composes a guided path across supported capabilities, offers partial/advisory support, or explains a genuine limitation. It must never invent an execution capability or promise an external outcome it cannot control. Once a valid path is identified, the product asks only for the minimum missing information, requests account verification at the point where work must be saved, and continues inside a newly created or existing EnterprateAI workspace.

The experience is chat-first but form-backed. Business-goal chips and quick-task chips are discovery accelerators, not separate products: both feed the same Goal Resolver, TaskSession and Assistant experience. Natural language establishes intent and context; structured inline controls capture dates, amounts, customer names, offerings, choices and other facts reliably. Confirmed information is written to the appropriate authoritative EnterprateAI records rather than remaining trapped in chat history. The Adaptive Dashboard is launched only after the Agent has prepared something useful, so the user’s first dashboard experience is an outcome - such as a funding-readiness assessment, invoice ready for review, or clear guided next-action plan - rather than an empty workspace.

North-star experience Business problem or task -> understood outcome -> validated capability path -> value first -> identity when needed -> progressive data capture -> useful output -> Adaptive Dashboard with work already prepared.

2.1 Product statement

Tell EnterprateAI what you want to achieve. EnterprateAI resolves the underlying business outcome, checks what it can genuinely support, works out what information it needs, preserves the goal/task, creates or connects the workspace, prepares supported work, and brings the user into the dashboard when there is something useful to review or act on. If the exact requested outcome is not directly executable, the Assistant should still guide the user toward the closest supported capabilities without misrepresenting what EnterprateAI can do.

2.2 Why this aligns with EnterprateAI

- The Assistant remains the conversational interface; Agentic Orchestration performs supported executable work behind the same experience.
- The Adaptive Dashboard remains the post-acquisition decision and action surface, not a separate onboarding product.
- Business, Customer, Offering, Invoice, Quotation and other operational entities remain authoritative records; chat is an interaction layer, not a second data store.
- Marketplace Claim remains a separate acquisition route, but both routes converge on the same user, Business, membership, entitlement and Adaptive Dashboard architecture.

3. Product Problem

A conventional SaaS signup asks a visitor to register, create a business, configure settings and learn navigation before receiving value. This is especially costly for micro and small business users who arrive with a specific job to complete. EnterprateAI also risks presenting a broad feature set before the user understands which capability is relevant.

The product problem is therefore to convert the visitor’s immediate business goal into a structured, safe and reusable task while collecting only the information required at that moment. The system must feel conversational and intelligent without sacrificing data quality, permissions, verification, deterministic calculations or approval controls.

4. Goals and Non-Goals

Goal	Requirement
Fast first acquisition	Let a visitor start useful work before a traditional setup flow.
Fast time to value	Produce a useful draft, assessment or plan as early as the selected task allows.
Progressive profiling	Ask only for facts required by the current task and reuse known data later.
Structured data capture	Translate confirmed conversation into canonical business records and task inputs.
Seamless AI experience	Use the existing Assistant + Agentic Orchestration model rather than a separate onboarding bot.
Reserve human authority	Draft and prepare freely where permitted; consequential external actions remain approval-controlled.

Protect GTM promise	Keep essential tool access frictionless and avoid consuming the user's 50 AI Credits for basic account setup and structured onboarding.
Converge acquisition routes	Homepage, Marketplace Claim and later referral/deep-link journeys converge on the same workspace architecture.
Successful goal resolution	Understand the underlying business outcome before declaring a request unsupported; map it to registered capabilities or provide an honest guided fallback without dead ends or fabricated execution.
Outcome-led discovery	Present the business owner's problems and desired outcomes more prominently than internal feature names, while retaining a quick-task route for users who know exactly what they want to create.
One resolver for every entry mode	Business-goal chips, quick-task chips and free-text requests all enter the same Goal Resolver, capability validation, TaskSession and workspace flow.

4.1 Non-goals

- A separate onboarding application or permanent onboarding dashboard.
- A pure free-text chatbot that stores important business facts only in conversation history.
- A universal questionnaire shown to every user regardless of task.
- Autonomous sending, signing, payment, funding submission or other consequential actions without the governing Agentic Orchestration approval policy.
- Replacing the existing Business Operations, Catalogue, readiness or Decision Intelligence modules.
- Inventing missing business, customer, tax, pricing, traction or financial facts.

5. Adopted Product Principles

Principle	Product application
Problem/task, not signup	The first prominent action is “What do you want EnterprateAI to do for your business?” rather than “Create account”.
Text-first, form-backed	Use natural language for intent and explanation; use structured fields/chips/selectors for facts where precision matters.
Only what is missing	Never ask the user to re-enter a fact already known and current unless confirmation is required.
Continue before dashboard	Do not launch an empty dashboard after verification. Continue the task until a useful draft/result is available or a clear blocking input is needed.
Draft first, commit second	Unverified or extracted values stay in the task draft until the user confirms them or governing policy allows persistence.
No fabricated completeness	When information is insufficient, state what is missing and continue progressively; do not invent values to finish the task.

source of truth	Confirmed customer, offering, business and operational facts are persisted in the appropriate authoritative record.
credits fund intelligence	Account creation, verification and basic structured onboarding are zero-credit; AI generation/analysis/automation follows the credit/entitlement policy.
outcome before feature name	Do not reject a request simply because the user did not name an EnterprateAI feature. First infer the underlying business outcome, then validate possible routes against the registered capability/task registry.
capability invention	The Assistant may explain and combine supported capabilities, but it must not claim EnterprateAI can execute, guarantee or automate an outcome unless that capability is registered and permitted.
business language before feature language	Lead with outcomes such as funding, customers, pricing, launch, cash flow, growth and risk. Use feature names as secondary quick tasks or as the supported route revealed by the Assistant.
no entry modes, one experience	A business-goal selection and a quick-task selection must never create separate onboarding architectures. Both resolve into the same Assistant, TaskSession, canonical Business and Adaptive Dashboard.
homepage suggestion integrity	Do not display a business-goal suggestion as a marketing promise unless at least one meaningful registered capability path exists. The user-facing copy must reflect the real execution boundary.

6. Homepage Entry Experience

The existing homepage remains a product marketing page, but the centred hero becomes the primary outcome-and-task entry surface. The visual hierarchy is deliberately outcome-led: business owners first see common problems and goals in their own language; users who already know the document or tool they need can use the smaller “Or start a task” row. Free-text input remains available at all times.

element	Requirement
row	AI-Native Business Decision Intelligence
headline	Build a More Resilient Business with Intelligence.
supporting line	Plan, operate, sell and grow from one intelligent business workspace.
primary prompt	What do you want EnterprateAI to do for your business?
secondary - Popular business goals	I need funding; I need more customers; Help me price my product/service; I want to launch a product/service; Improve my cash flow; Help me grow my business; Test a business idea; Help me understand my business risks.
assurance	Free essential business tools; 50 AI Credits to get started.
secondary - Or start a task	Create Invoice; Create Quotation; Create Business Plan; Prepare Proposal. Additional quick tasks may be enabled only when they have a valid registered path.

al hierarchy

gestion governance

Popular business goals are the primary, more prominent chips. Quick tasks are visually secondary. Neither replaces the natural-language input.

Goal and task suggestions are configuration/registry-driven. A displayed suggestion must resolve to a current direct, composite or partial supported path; its copy must not imply a guarantee such as “Get funding” or “Get customers”.

UX rule The homepage supports three equivalent inputs into one intelligence path: choose a popular business goal, choose a quick task, or type a free-text request. Business goals are visually primary because they match how owners describe needs; quick tasks remain available for speed. Every selection enters the same Goal Resolver and capability validation flow. Suggestions accelerate discovery but never restrict natural-language input.

6.1 Default homepage business-goal routing

Homepage goal	Intended outcome route (subject to registered capability)	First useful response / boundary
I need funding	Funding Readiness + relevant Business Plan/financial evidence workflows	Assess readiness, requirement, use of funds and gaps; never promise funding or investor approval.
I need more customers	Growth intelligence + Marketplace visibility/opportunity + Proposal/Quotation support	Help improve the route to customers; never promise or directly guarantee customers.
Help me price my product/service	Viability/economics + offering/pricing data + scenario support where registered	Help model price, cost, margin and break-even implications; user confirms commercial price.
I want to launch a product/service	Launch Readiness + Idea Validation + scenario/planning support	Assess launch readiness, gaps, assumptions and next actions; do not fabricate market evidence.
Improve my cash flow	Survival/financial intelligence + scenario + receivables/expense workflows where registered	Identify cash drivers, risks and possible actions; calculations remain deterministic where applicable.
Help me grow my business	Growth Engine + capacity/market/commercial data + scenario support	Identify growth constraints/opportunities and a supported action path; no guaranteed growth result.
Test a business idea	Idea Validation / Service Idea Viability	Collect minimum economics/demand/delivery inputs and produce explainable viability output.
Help me understand my business risks	Fragility + Survival + Stability intelligence as available	Surface explainable risk signals and next actions from authoritative engines/data.

These mappings define intended product behaviour, not hard-coded promises. During implementation, each homepage goal must be validated against the current production capability registry and owning PRD before it is enabled.

7. Universal End-to-End Journey

Visitor lands on homepage hero
↓
Chooses Popular Business Goal / Quick Task / submits free-text request
↓
Assistant resolves underlying business outcome -> validates current registered capabilities -> classifies direct / composite / partial / no-match / blocked path
↓
Assistant explains the supported route in business language and asks 1-3 high-value questions only when they improve routing or the selected workflow
↓
System reaches persistence checkpoint when work must be saved/continued reliably
↓
Request email -> verify email -> create/authenticate user account
↓
Collect name + business name if known/needed; preserve original goal, selected entry mode and pre-auth context
↓
Resume the goal/task; load authorised known business data and ask only missing required inputs
↓
Confirm proposed records / task payload / goal plan where applicable
↓
Persist authorised canonical records + create supported task/workflow state
↓
Agent / engine / tool prepares a useful output or next-action plan
↓
Launch existing Adaptive Dashboard with active goal/task/result visible
↓
User reviews -> approves/continues/sends/reassesses according to capability policy

7.1 Persistence checkpoint

The system may collect a small amount of low-risk task information before authentication in an ephemeral task session. It must request email verification before creating persistent private business records, documents or agent workflows tied to the user. The exact anonymous-session retention period is a configurable security/legal decision and must be disclosed.

7.2 Identity sequence

1. Ask for email only when the system needs to save or continue the task reliably.
2. Verify email using the current authentication mechanism.
3. Create or authenticate the user account.
4. Ask for the user's name; business name may remain optional until a task or canonical business record requires it.
5. Resume the same TaskSession with no loss of previously supplied information.
6. Resolve an existing Business/membership before creating a new Business record.

8. Interaction Model: Chat-First, Form-Backed

Interaction type	Use	Example
Conversation	Understand goals, explain progress, ask open questions, clarify ambiguity, summarise what is known.	"Create an invoice for ABC Consulting."
Suggestion chips	Speed common choices without typing.	7 days / 14 days / 30 days / Custom
Inline structured fields	Capture facts requiring precision and validation.	Amount, date, quantity, currency, customer email
Review cards	Show interpreted facts before persistence or consequential action.	Customer: ABC Consulting Ltd; Service: Website redesign; Amount: £1,500
Progressive checklist	Show useful completion state for larger tasks without exposing a long questionnaire.	Business Plan 62% ready; 4 required items remain

8.1 Task Definition contract

Every supported homepage task should be represented by a versioned TaskDefinition or equivalent configuration, mapped onto existing production services rather than implemented as bespoke conversational logic.

Field	Purpose
task_key	Stable identifier, e.g. create_invoice, create_quotation, business_plan, idea_validation.
supported_intents	Natural-language goals that may resolve to the task.
required_inputs	Facts required to reach the task completion condition.
optional_inputs	Facts that improve quality but must not block basic completion.

authoritative_sources	Existing records to read before asking the user.
persistence_targets	Canonical records to create/update only after confirmation/permission.
completion_condition	Minimum state at which a useful output can be produced.
workflow_capability	Agentic Orchestration capability/tool mapping where execution is supported.
credit_policy	Zero-credit, included first-task allowance, or credit-consuming AI work.
approval_policy	Actions requiring explicit review/approval.

Critical rule The LLM may infer intent and propose structured fields, but required commercial or business facts must be confirmed or retrieved from authorised sources before they become authoritative.

8.2 Goal Resolution, Capability Mapping and Graceful Fallback

Users will often describe a desired business outcome rather than a product feature. Examples include “I need more customers”, “I need to improve cash flow”, “help me grow”, or “I need to win more work”. The onboarding layer must treat these as goal-resolution problems, not immediate unsupported-feature errors. It must first understand the underlying outcome, then validate possible routes against registered EnterprateAI TaskDefinitions, tools, engines and product capabilities.

Resolution class	Definition	Required behaviour
Direct supported task	The goal maps confidently to one registered TaskDefinition/capability.	Start that task, reuse known data, and ask only missing inputs.
Composite business goal	The user asks for an outcome that is not a single feature but can be advanced through a combination of supported capabilities.	Explain how EnterprateAI can help, identify the most relevant supported capabilities, ask the minimum clarifying question(s), and create a guided path. Do not promise the business outcome itself.
Partial / advisory support	EnterprateAI can analyse, plan, recommend or prepare supporting work, but cannot execute the user’s requested real-world outcome directly.	State the execution boundary clearly, then offer the supported analysis, planning, Marketplace, document or operational path that moves the user toward the goal.
No valid capability match	No registered EnterprateAI capability materially addresses the request.	Say that EnterprateAI does not currently perform that action. Offer relevant product navigation or planning help only where genuinely useful; otherwise stop cleanly without inventing functionality.

Blocked by permission / entitlement

A capability exists but is unavailable to this user/business because of role, plan, credits or approval rules.

Explain the access constraint, preserve the goal/task, and offer the authorised next step such as approval, upgrade or zero-credit alternative.

8.2.1 Goal Resolver contract

Field	Purpose
original_goal	Exact user request as submitted.
underlying_outcome	Normalised business outcome/problem the user is trying to achieve.
resolution_class	direct_supported composite_goal partial_support no_match blocked.
candidate_capabilities	Ranked registered TaskDefinitions/tools/engines that materially address the outcome.
confidence	Resolution confidence used to decide whether one clarifying question is required before routing.
execution_boundary	What EnterprateAI can prepare/analyse/execute and what remains outside its control.
recommended_path	Ordered supported next actions presented to the user.
credit_implication	Whether the next selected action is zero-credit or consumes AI Credits.
provenance	Registry/rule/version used to validate the capability mapping.

8.2.2 Mandatory resolution rules

- Do not classify a goal as unsupported merely because the wording does not match a feature or menu label.
- The LLM may infer the user's underlying outcome, but it may recommend execution only from the current registered capability/task/tool catalogue or equivalent production source of truth.
- If confidence is low or two materially different routes are plausible, ask one concise clarifying question before selecting a path.
- For composite goals, present a short "Here is how I can help" path using supported capabilities; do not expose a long internal feature list.
- Never guarantee external outcomes such as getting customers, winning contracts, obtaining funding, increasing revenue or achieving a specific growth result.
- Basic goal interpretation, capability discovery, limitation explanation and routing are zero-credit onboarding behaviour. Credits begin only when the user elects a substantive AI generation, analysis, simulation or automation action governed by the credit policy.
- A no-match response must remain useful but truthful: explain the current limitation, preserve any relevant context, and offer planning/navigation help only when it has genuine value.

8.2.3 Example - "I need more customers"

Step	Expected EnterprateAI behaviour
Interpret	Resolve the request as a customer-acquisition / growth outcome, not as a literal feature name.
Validate capabilities	Check registered capabilities such as Growth intelligence/assessment, Marketplace visibility/opportunity settings, Proposal/Quotation preparation and relevant business/offer/customer data workflows.

Set boundary	Explain that EnterprateAI can help the user work toward acquiring customers but cannot guarantee or directly deliver customers.
Clarify	Ask only the most useful questions, for example: What do you sell? Who is your ideal customer? Where do you currently get customers? Reuse existing Business/Offering/Customer data if already known.
Route	Start the most relevant supported assessment/workflow or compose a short action path across supported capabilities.
Persist	Before authentication, keep answers in the ephemeral TaskSession. After verification, confirm before writing reusable business/market/customer facts to canonical records.
Credits	Goal interpretation and routing are zero-credit. Any subsequent AI-heavy growth analysis, generation or automation follows its registered credit policy.

Illustrative user-facing response: “I can help you work toward getting more customers. I can assess your current offer and target market, identify growth gaps, improve how your business is positioned in the Marketplace, and help prepare proposals or quotations. I cannot guarantee or directly acquire customers for you. To start, what do you sell and who are you trying to reach?”

9. Initial Task Journeys

9.1 Create Invoice

Area	Requirement
Initial prompt	“Create an invoice for ABC Consulting Ltd.”
Minimum early questions	Customer; service/item; amount or quantity/price where not inferable.
After authentication	Reuse existing Customer, Offering, BusinessProfile, invoice settings and pricing where available.
Likely missing fields	Seller identity, customer destination, currency, due date/payment terms, tax/VAT treatment where applicable.
Persistence	Customer -> Customer Catalogue; service/item -> Offering/Catalogue; confirmed commercial inputs -> draft Invoice.
Completion condition	Valid draft invoice can be rendered for review.
Dashboard handoff	Agent card: “Invoice for ABC Consulting Ltd is ready for approval.”
External effect	Sending requires the applicable approval/permission policy.

9.2 Create Sales Quotation

Area	Requirement
Core inputs	Customer, offering/scope, quantity, price, validity, commercial terms where required.
Reuse	Customer + Catalogue + approved pricing + business settings.
Completion condition	Reviewable quotation draft.
Handoff	Adaptive Dashboard / Agent shows quotation ready for review; sending follows Agentic approval policy.

9.3 Business Plan Generator

Area	Requirement
Core behaviour	Progressive multi-turn collection; never present a giant upfront form.
Reuse	BusinessProfile, Offerings, target customer, Idea Validation, financial assumptions, Launch/Funding data where authorised and current.
Required gaps	Ask one logical cluster at a time - business/offer, customer/problem, model/pricing, delivery, market/GTM, economics/forecast as applicable.
Partial result	Generate/preview completed sections while clearly showing missing sections and assumptions.
Completion state	Example: “Business Plan 62% ready - 4 required items remain.” Exact completeness logic must be deterministic/configured, not an LLM guess.
Handoff	Dashboard surfaces “Continue Business Plan” or completed draft with next useful action.

9.4 Idea Validation

Area	Requirement
Core inputs	Idea/service description, target customer, pricing/revenue assumptions, direct/fixed cost inputs, capacity and demand evidence as required by the validated engine profile.
Behaviour	Conversational collection + structured numeric fields; explain why each required input matters.
Authority	Deterministic Viability Engine owns calculations and outcome; LLM explains and asks for missing inputs.
Handoff	Dashboard shows result, risk flags and next action; no fabricated score when required data is missing.

9.5 “I need funding” - Funding Readiness Journey

“I need funding” is a primary outcome-led homepage route because the user is expressing a business need, not asking for a feature. EnterprateAI must translate the request into a Funding Readiness journey: understand the funding requirement, assess how prepared the business is, identify evidence/material gaps, recommend next actions and prepare supported materials. EnterprateAI must never imply that it can guarantee funding, investor approval or lender acceptance.

Area	Version 2 requirement
Homepage trigger	User selects “I need funding” or expresses an equivalent funding need in natural language.
Initial response	Explain that EnterprateAI can help assess funding readiness, identify gaps, organise evidence and prepare supported materials - not guarantee funding.
Minimum early questions	Funding amount sought; intended use of funds; business stage. Ask only what is not already known.
Persistence checkpoint	After enough context exists to start the assessment, request email verification so progress can be saved and the workspace created/linked.
Post-auth reuse	Reuse authorised BusinessProfile, financials, customers/traction, Offering, Business Plan, Launch Readiness and prior evidence before asking the user again.
Progressive evidence collection	Collect only missing readiness inputs in logical clusters: funding need/use, business model, traction, financial position/forecast, growth case, risks/gaps and supporting evidence.
Useful first output	A Funding Readiness result/score or equivalent registered assessment, strengths, material gaps and prioritised next actions. Deterministic scoring rules remain authoritative where defined.
Dashboard handoff	Open the Adaptive Dashboard with “Funding Readiness” active: current readiness, missing evidence/materials, recommended next actions and continuation path visible.
Related supported actions	Where available and entitled: continue Business Plan, prepare investor/funding materials, improve forecasts/evidence, run scenarios, reassess readiness.
Credit treatment	Goal interpretation, basic onboarding and structured data collection are zero-credit. Substantive AI generation/analysis follows the Funding Readiness/feature credit policy.
Boundary	Never say “we will get you funding”. Use wording such as “I can help you become more funding-ready and prepare what you need.”

Other outcome-led suggestions and quick tasks use the same universal contract: resolve intent, validate the current capability path, reuse known data, ask only missing inputs, preserve human authority and hand off to the existing Adaptive Dashboard when useful work is available.

10. Data, Ontology and Source-of-Truth Rules

Conversation is not the source of truth. The onboarding layer must maintain an explicit separation between temporary conversation/task context and authoritative business records.

Information	Ownership	Direction
TaskSession	Temporary onboarding/task context: original goal, resolved task, pre-auth inputs, authentication state, missing-input state, source channel and correlation ID.	New or mapped to existing session/workflow infrastructure.
User / Account	Authenticated person and account identity.	Existing authentication/account system.
Business / BusinessProfile	Canonical business identity, stage and context.	Existing ontology/system of record.
Customer	Confirmed customer identity/contact.	Catalogue / Customer record.
Offering	Confirmed product/service and approved pricing where supported.	Catalogue / Offering record.
Invoice / Quotation / Contract / Receipt	Authoritative commercial records and states.	Business Operations.
Business Plan / Proposal documents	Owned by the relevant document/blueprint feature.	Business Blueprints / document service.
Assessments / Engine outputs	Authoritative deterministic results.	Assessment/engine services.
WorkflowRun / Approval	Execution status, approvals and audit.	Agentic Orchestration.

10.1 Write policy

- Pre-auth values stay in TaskSession only; do not create private Customer, Offering, Business Operations or assessment records for an anonymous visitor.
- After authentication, proposed extracted values remain draft until confirmation or an explicitly safe deterministic rule permits persistence.
- When a matching Customer or Offering already exists, prefer reuse/confirmation to duplicate creation.
- Never silently overwrite canonical business facts from conversational text. Material updates require user confirmation and the current write permission.
- Preserve provenance: user-entered, inferred/proposed, existing record, integration, assessment or other authorised source.

10.2 Duplicate prevention

Before creating a Business, Customer, Offering or other reusable entity, the system must search/match within the authorised scope and apply a deterministic duplicate-resolution policy. Ambiguous matches require confirmation. Cross-tenant/private matches must never expose another business's data.

11. Onboarding / Task State Model

State	Meaning
initiated	Visitor has submitted a goal or selected a suggestion.
collecting_context	System is collecting minimal early task information.
awaiting_identity	Persistence checkpoint reached; email/authentication required.
awaiting_verification	Email verification is pending.
resuming	Authenticated session is reconnecting to the preserved task.
collecting_required_data	System is asking only remaining task-required inputs.
ready_to_prepare	Completion condition for draft/assessment preparation is met.
preparing	Agent/tool/engine is producing the first useful output.
awaiting_review	Output is ready for user review/approval.
handed_off	Adaptive Dashboard/module handoff completed.
blocked	Missing input, permission, entitlement or unsupported capability prevents continuation.
abandoned	Session expired/cancelled without conversion; retention policy applies.

12. Assistant, Agent and Tool Boundaries

Layer	Responsibility
Assistant / LLM	Understand intent, ask/clarify, extract proposed fields with provenance, explain progress, summarise known/missing information.
Task/Onboarding service	Track state, required fields, source records, authentication checkpoint, completion condition and handoff.
Deterministic engines	Calculate viability/readiness/simulation metrics and classifications.
Agentic Orchestration	Plan and execute supported workflows through registered tools, approvals, permissions, audit and recovery.

Business Operations / modules	Own authoritative commercial/document/assessment records.
Adaptive Dashboard	Surface active task, prepared output, approvals, next action and broader business priorities.

Agent rule The Assistant can make the experience feel conversational, but it cannot invent a capability, override an engine score, expand permissions or bypass the approval policy. Any claimed executable capability must be validated against the current registered TaskDefinition/tool/engine capability source of truth before it is offered as an action.

13. Free Tools, AI Credits and First-Task Economics

The onboarding experience must support EnterprateAI's GTM promise that essential business tools are genuinely accessible while preserving the 50 free AI Credits for intelligence and automation value.

Goal interpretation, capability matching, limitation explanation, clarifying-for-routing and fallback presentation are onboarding/navigation behaviours and must not consume the user's 50 AI Credits. Credit metering begins only when the user proceeds into a registered AI generation, analysis, simulation or automation action according to that capability's policy.

Activity	User credit treatment	Rationale
Email verification / account creation	0 user credits	Platform access and identity, not AI value.
Basic name/business/customer/offering data capture	0 user credits	Structured onboarding and data entry.
Deterministic form validation / record lookup	0 user credits	Core application behaviour.
Create/edit basic invoice, quotation, receipt and catalogue records	0 user credits where the action is deterministic/basic	Protect free essential-tool GTM promise.
Intent classification / lightweight onboarding assistance	Prefer 0 user credits; meter internally if needed	Onboarding friction should not burn the user's allowance.
AI narrative generation, complex drafting, extraction or reasoning	Credits apply according to pricing policy	Actual AI value.
Business Plan / Proposal generation	Credits apply, with any adopted first-task trial allowance	Higher-cost generative work.
Scenario / Decision Intelligence / advanced Agent workflows	Credits/plan entitlement applies	Premium intelligence/automation.
External sends / workflow actions	Credit and entitlement policy separate from approval	Action cost does not remove approval requirements.

Recommended commercial rule Give every new verified user one complete essential-tool first-task experience without consuming their 50 AI Credits. The credits should feel like extra intelligence, not the price of getting through onboarding.

14. Adaptive Dashboard Handoff

The homepage onboarding layer does not own a separate dashboard. Once the first useful output exists, the user enters the existing Adaptive Dashboard in the appropriate business context.

Pathway	Handoff behaviour
Startup	First-task output plus next actions, launch/viability/cash priorities according to existing dashboard logic.
Small Business	First-task output plus operational/risk/recommendation priorities.
Multi-Entity	Current claimed/created business opens in its entity scope; portfolio view remains governed by the existing account-scope model.

- The onboarding task may influence “Current Goal” priority but must not override stage, scope, role or subscription dimensions.
- A task such as invoice creation may surface in the Agent panel and Needs Approval area; the dashboard must not become a permanent onboarding wizard.
- If preparation cannot complete because required information is still missing, the dashboard may still be entered only when there is a useful persistent state to show; otherwise continue the focused task flow.

15. Functional Requirements

	Requirement	Priority
	Homepage hero exposes one central natural-language prompt plus two suggestion levels: prominent business goals/problems and secondary quick-task chips.	MVP
	A visitor can start a supported task without first creating an account.	MVP
	The system creates a TaskSession and preserves original goal, source and task context.	MVP
	Intent resolution identifies the underlying business outcome and maps any executable route only to registered/supported task definitions, tools or engines.	MVP
	The system asks only minimum early questions before the persistence checkpoint.	MVP
	Pre-auth information remains ephemeral and does not create canonical private business records.	MVP
	The task context survives email verification, signup/login, redirect, retry and resume.	MVP

	After authentication, the system searches authorised existing Business/Customer/Offering records before creating duplicates.	MVP
	The system dynamically asks only remaining required task inputs, using chat, chips or structured fields as appropriate.	MVP
	Required missing facts are never fabricated; the user receives a clear missing-information path.	MVP
	Confirmed inputs can write to canonical records according to ownership, permission and provenance rules.	MVP
	Invoice and quotation tasks can produce reviewable drafts from valid authorised data.	MVP
	Business Plan and Idea Validation can operate progressively and preserve partial progress.	MVP
	Deterministic engines own calculated scores/classifications; onboarding does not duplicate scoring.	MVP
	The first useful output is surfaced before/at Adaptive Dashboard handoff, not after an empty dashboard launch.	MVP
	Effectful actions route through existing Agentic Orchestration approvals/permissions.	MVP
	Zero-credit onboarding/basic-tool actions are distinguished from credit-consuming AI activities.	MVP
	Free-credit exhaustion never deletes task/business records or blocks zero-credit essential-tool use where entitled.	MVP
	Marketplace Claim users and homepage users converge on the same canonical account/business/dashboard model.	MVP
	Analytics capture funnel progress without logging raw sensitive task content.	MVP
	Task definitions can be versioned and expanded without redesigning the homepage onboarding framework.	MVP
	A request that does not directly match a feature is first resolved to its underlying business outcome. The system then maps it to direct, composite or partial supported capabilities; only a true no-match is treated as unsupported. Limitations are explicit and execution is never invented.	MVP
	Future voice input can resolve into the same TaskSession/intent contract.	Later-ready
	The Goal Resolver classifies each goal as direct_supported, composite_goal, partial_support, no_match or blocked and records the validated capability mapping/provenance.	MVP
	Composite business goals can produce a concise guided path across multiple registered EnterprateAI capabilities without presenting those capabilities as a guaranteed outcome.	MVP
	Low-confidence or materially ambiguous capability mapping asks one concise clarifying question before routing; high-confidence mappings proceed without unnecessary questioning.	MVP

Goal interpretation, capability discovery and fallback routing are zero-credit; subsequent credit-consuming activity follows the selected capability policy.	MVP
Capability suggestions are sourced from the current registered capability/task/tool catalogue or equivalent production source of truth; the LLM cannot create a capability from model knowledge alone.	MVP
Homepage prompt is “What do you want EnterprateAI to do for your business?” and accepts free-text business goals as well as feature/task language.	MVP
The homepage presents Popular Business Goals as the primary suggestion group and “Or start a task” as a visually secondary quick-task group.	MVP
Default business-goal suggestions include funding, more customers, pricing, launch, cash flow, growth, idea testing and business risk, subject to current registered capability validation.	MVP
Every displayed homepage goal must resolve to a meaningful registered direct/composite/partial path; suggestions with no valid current path are not displayed as marketing promises.	MVP
Selecting “I need funding” routes into the Funding Readiness journey, preserves the user’s funding goal through verification, and never promises funding approval or access.	MVP
Business-goal chips, quick-task chips and free-text requests all enter the same Goal Resolver, TaskSession, capability validation and canonical workspace flow.	MVP
Homepage suggestion labels/order can be changed through controlled configuration or equivalent production mechanism without changing core onboarding architecture.	MVP
Analytics records entry_mode and selected_goal/task key so product can measure which owner problems drive activation, first value, retention and conversion.	MVP

16. UX Screen / Surface Inventory

	Surface	Purpose
	Homepage Hero Task Launcher	Centred natural-language prompt; primary Popular Business Goals; secondary Quick Tasks; reassurance about free tools/credits.
	Conversational Task Panel	Assistant response plus dynamic inline fields/chips.
	Email Persistence Checkpoint	Explain that email is needed to save work; collect email.
	Email Verification	Code/link verification using current authentication.
	Identity Basics	Name + business name if known; avoid unrelated business questionnaire.
	Task Resume	Show preserved goal and known information after authentication.
	Adaptive Missing-Data Collection	Task-specific structured questions only.
	Review Known Information	Confirm proposed Customer/Offering/business/task data before persistence where required.

	Preparing State	Agent/tool/engine progress with concise, truthful status.
	Adaptive Dashboard Handoff	Prepared output visible in Agent/approval/task surface.
	Goal Path Preview	After a business goal is selected, explain “Here is how I can help” in plain business language, show the first relevant question(s), and state any external-outcome boundary before signup.
	Funding Readiness First-Value State	For “I need funding”, show funding requirement/use/stage collection, assessment progress and readiness/gap output before/at dashboard handoff.

17. Required UX States and Recovery

State	Required behaviour
High-level / non-feature goal	Resolve the underlying business outcome, validate direct/composite/partial capability matches, explain the supported path and ask only the minimum clarifying question(s). Do not reject merely because the goal is not a feature name.
Ambiguous intent	Ask one clarifying question before starting a task definition.
Email already registered	Authenticate and resume the TaskSession rather than creating another account.
Verification failed/expired	Retry or resend without losing task context.
Potential duplicate Business	Ask user to select/join authorised business or enter a controlled resolution path.
Missing required input	Explain exactly what is missing and why; preserve all completed progress.
LLM unavailable	Keep structured inputs and deterministic application functionality available; provide retry/manual path.
Tool/engine failure	Isolate failure, preserve TaskSession and canonical records, show safe retry state.
Credits exhausted	Continue zero-credit eligible operations; explain which AI action is restricted and offer upgrade.
Permission restricted	Explain that the user cannot perform the action; do not leak inaccessible data.
Partial-support request	Explain exactly what EnterprateAI can analyse, prepare or recommend and what it cannot execute; offer the closest supported path without guaranteeing the outcome.
No-capability match	State that EnterprateAI does not currently perform the requested action. Offer relevant planning/navigation help only if genuinely useful; otherwise end cleanly without invented execution.
Homepage goal selected	Respond in the language of the business outcome, not by dumping an internal feature list. Explain the nearest supported path and ask the minimum useful next question.
Funding goal selected	Explain readiness support and the no-guarantee boundary, then collect funding

amount/use/business stage before the persistence checkpoint as appropriate.

18. Conceptual API and Event Contracts

Names are conceptual and must be mapped to the current production stack during Claude Code production discovery. The PRD defines behaviour, not a requirement to create duplicate services.

Conceptual operation	Purpose
POST /task-sessions	Start task from homepage goal; return session ID and resolved/clarification state.
GET /task-sessions/{id}	Read current task, known/missing fields, auth state and next interaction.
POST /task-sessions/{id}/inputs	Submit structured/chat-derived confirmed inputs.
POST /task-sessions/{id}/attach-user	Bind authenticated user after verification.
POST /task-sessions/{id}/resume	Resume after auth/redirect/retry.
POST /task-sessions/{id}/prepare	Start eligible draft/assessment/workflow when completion condition is met.
GET /task-sessions/{id}/status	Read preparation/workflow status and handoff destination.
POST /goal-resolution (conceptual)	Resolve original goal -> underlying outcome -> resolution class -> validated capability candidates -> execution boundary -> recommended path. May be implemented inside TaskSession service rather than as a separate endpoint.

18.1 Events

- HomepageGoalSubmitted
- TaskSessionStarted / TaskIntentResolved / TaskClarificationRequested
- AuthenticationRequired / EmailVerified / TaskSessionAttachedToUser
- TaskInputConfirmed / CanonicalRecordLinkedOrCreated
- TaskReadyToPrepare / FirstOutputPrepared / DashboardHandoffCompleted
- TaskBlocked / TaskFailed / TaskAbandoned
- Existing Agentic events continue for workflow/approval/tool execution; do not duplicate them in onboarding.
- GoalOutcomeResolved / CapabilityMatchClassified
- CompositeGoalPathPresented / FallbackPathAccepted
- NoCapabilityMatch

19. Security, Privacy and Trust Requirements

- Treat all pre-auth chat/task content as untrusted input; it cannot change system policy or tool permissions.
- Do not create private persistent business/customer records until the user is authenticated and the relevant write policy permits it.
- Preserve minimum necessary pre-auth data; exact retention and deletion policy must be reviewed before launch.

- Do not expose private existing Business/Customer/Offering matches to an unauthorised user during duplicate resolution.
- Every canonical write and consequential workflow preserves actor, source, timestamp and correlation/audit references.
- Sensitive values must not be copied into analytics payloads or routine logs.
- Account/business scope, role and entitlement are rechecked at the service boundary before tool execution.
- Voice, future integrations or deep links never bypass verification or approval requirements.

20. Analytics and Success Measures

homepage_goal_prompt_viewed
↓
entry_mode_selected (business_goal quick_task free_text)
↓
goal_or_task_submitted
↓
goal_outcome_resolved + capability_match_classified
↓
first_useful_question_answered
↓
verification_started
↓
email_verified
↓
task_or_goal_path_resumed
↓
required_inputs_completed
↓
first_output_prepared
↓
dashboard_handoff_completed
↓
first_output_approved_or_completed
↓

Metric	Purpose
Primary activation	Verified user who completes or reaches review-ready state on the first chosen task.
Time to first useful interaction	Time from goal submission to a meaningful task-specific response.
Time to first output	Time from goal submission to reviewable draft/assessment/result.
Pre-auth -> verified conversion	How many task starters verify email.
Verified -> first output conversion	Whether onboarding actually delivers value.
Data reuse rate	Percentage of required fields fulfilled from existing authorised records rather than re-entry.
Question burden	Median user inputs/questions required by task type.
First-task completion	Invoice/quote/etc. completion or review-ready rate.
Credit preservation	Credits remaining at first dashboard handoff for essential-tool tasks.
Quality	Draft correction rate, intent correction rate, duplicate-record correction rate.
Retention/conversion	7-day/30-day return and paid conversion by first-task type.
Goal resolution success	Percentage of submitted goals that resolve to a direct, composite or partial supported path without a dead end.
Fallback acceptance	Percentage of composite/partial/no-direct-feature users who accept a recommended supported next step.
True no-match rate	Percentage of goals with no valid registered capability match; used to identify product demand without fabricating capability.
Capability hallucination rate	Must remain zero: suggested executable capabilities not present in the registered source of truth.
Business-goal engagement	CTR and submission rate for each Popular Business Goal compared with Quick Tasks and free-text entry.
Goal-to-first-value conversion	For each business goal, percentage reaching a meaningful assessment, draft, recommendation or guided plan after verification.
Funding funnel	I-need-funding selection -> verification -> readiness assessment started -> readiness output -> continuation/paid conversion.

21. MVP Scope and Delivery Sequence

Phase	Deliverable	Exit direction
Phase 0	Production mapping	Audit current homepage, auth, onboarding, Business/Customer/Offering creation, invoice/quotation flows, dashboard routing, Agent runtime, credits and analytics. No coding.
Phase 1	Outcome-led homepage launcher	Centred chat-first input; primary Popular Business Goals; secondary Quick Tasks; ephemeral TaskSession; shared Goal Resolver and capability validation.
Phase 2	Identity checkpoint	Email verification, account attach/resume, context preservation and duplicate-safe business resolution.
Phase 3	Invoice first-task technical pilot	Prove persistence, data reuse, canonical writes, draft creation and dashboard/approval handoff on a structured low-ambiguity flow.
Phase 4	Funding outcome pilot	Enable “I need funding” from homepage through Funding Readiness, progressive evidence collection and outcome-led dashboard continuation.
Phase 5	Quotation + Idea Validation	Reuse the same task/goal framework for quotation and idea-testing entry paths.
Phase 6	Business Plan + Launch Readiness	Progressive multi-section planning and product/service launch outcome path.
Phase 7	Growth / customer / cash / risk goal expansion	Enable additional primary business-goal chips only where current registered capability paths pass product acceptance tests.
Phase 8	Optimisation	Analytics-driven goal ordering, contextual suggestions for signed-in users, smarter reuse, voice-ready surface and funnel experiments.

Pilot recommendation Use Invoice first to validate the technical onboarding/orchestration plumbing because it is structured and easy to test. In parallel, make “I need funding” the first high-value outcome-led acquisition pilot once Funding Readiness integration is production-mapped, because it demonstrates the stronger problem-to-intelligence proposition of EnterprateAI.

21.1 MVP exclusions

- Unlimited open-ended autonomous task execution.
- Voice input launch itself; architecture should remain voice-ready.
- Predictive personalisation of task suggestions before baseline funnel data exists.
- Complex external integrations not required for the first invoice/quotation journey.
- Replacing the Marketplace Claim flow; it remains a separate acquisition route that converges downstream.

22. Acceptance Criteria

Acceptance criterion

	A new visitor can submit “Create an invoice” from the homepage without registering first.
	The system preserves the original goal and pre-auth task inputs through email verification.
	No private canonical Customer/Offering/Invoice record is created for an anonymous visitor.
	An existing registered email signs in and resumes the same task instead of creating a duplicate account.
	After authentication, existing authorised Customer/Offering records are reused or confirmed before duplicate cre
	The invoice flow asks only fields still required after existing data is loaded.
	The system never invents missing customer, price, tax, payment term or business identity data.
	A valid invoice flow can produce a reviewable invoice draft and hand the user into the Adaptive Dashboard with th visible.
	The user does not land on an empty dashboard immediately after email verification.
	Sending the invoice follows current permissions and Agentic Orchestration approval policy.
	Email verification, account creation and basic onboarding do not consume the user’s 50 AI Credits.
	Zero-credit essential-tool operations remain usable when AI Credits are exhausted, subject to entitlement.
	Business Plan flow preserves partial progress and identifies missing required sections without fabricating comple
	Idea Validation outputs come from the deterministic engine; LLM failure cannot fabricate or replace the score.
	TaskSession state is recoverable after refresh, redirect or transient failure according to retention policy.
	Cross-tenant matching never exposes private business/customer data.
	Homepage acquisition and Marketplace Claim acquisition converge on the same canonical Business/membersh architecture.
	Analytics record funnel states without raw document bodies, sensitive financial content or private customer deta
	A non-feature, unsupported or ambiguous goal follows the Goal Resolver path: underlying outcome -> capability validation -> clarification where needed -> direct/composite/partial/no-match response. N execution is presented.
	Existing direct login, Marketplace, invoice, quotation and dashboard journeys continue to function during staged

- When a new visitor enters “I need more customers”, the system does not reject the request solely because there is no “Get Customers” feature. It resolves the growth/customer-acquisition outcome and presents only the registered capabilities.
- The “I need more customers” path explicitly states that EnterprateAI can help the user work toward that goal but cannot guarantee or directly deliver customers.
- A capability or executable action not present in the registered TaskDefinition/tool/engine source of truth is presented as something EnterprateAI can perform.
- Goal resolution and fallback presentation do not consume AI Credits; any later credit-consuming action is metered only according to the selected capability policy.
- A true no-match produces an honest limitation response and useful navigation/planning option when appropriate, without forcing signup or creating a fake task.
- Homepage main prompt reads “What do you want EnterprateAI to do for your business?” and remains open to free text.
- Popular Business Goals are displayed more prominently than Quick Tasks, with “I need funding” and approved goal suggestions available only when they have a valid registered support path.
- Quick Tasks remain visible as a secondary “Or start a task” route and include at minimum Create Invoice, Quotation, Create Business Plan and Prepare Proposal when enabled/entitled.
- Selecting a goal chip, task chip or typing equivalent free text produces the same normalised Goal Resolver/TaskSession behaviour rather than separate onboarding implementations.
- Selecting “I need funding” produces a funding-readiness response, asks only relevant early inputs, provides context through verification, and never states or implies that EnterprateAI guarantees funding.
- After the funding path has sufficient authorised data, the user receives a meaningful readiness/gap/next output and reaches the Adaptive Dashboard with that funding context visible.
- No homepage business-goal chip is enabled if production capability validation returns a true no-match.
- Analytics distinguish business_goal, quick_task and free_text entry modes and record the selected goal without storing sensitive raw content unnecessarily.

23. Production Dependencies

- Current homepage/hero implementation and routing.
- Authentication, email verification, account and membership model.
- Canonical Business/BusinessProfile creation and stage/context rules.
- Customer and Offering/Catalogue APIs and duplicate matching behaviour.
- Invoice and quotation draft/create/render/send mechanisms.
- Adaptive Dashboard context/composition endpoint and Agent panel/approval integration.
- Agentic Orchestration Task/Workflow APIs, registered tools and approval model.
- AI Credit metering and entitlement services.
- Business Plan, Idea Validation, Proposal, Readiness and engine interfaces.
- Analytics, events, audit, notification, background job and recovery infrastructure.

24. Open Product Decisions Before Formal Adoption

Decision	Recommended direction / question	Owner
----------	----------------------------------	-------

auth session retention	Define disclosed retention duration and deletion behaviour.	Product + Security/Legal
act first-task credit allowance	Confirm zero-credit essential first task versus specific AI actions that consume credits.	Product + Commercial
business name optionality	Map when a provisional/minimal Business can exist without a formal business name and what tasks require seller identity.	Product + Engineering
homepage business goals at launch	Recommended Version 2 set: I need funding; I need more customers; Help me price my product/service; I want to launch a product/service; Improve my cash flow; Help me grow my business; Test a business idea; Help me understand my business risks. Enable only after production capability validation.	Product
invoice tax/VAT questions	Map to current production invoice rules; onboarding must not invent tax status.	Product + Finance/Engineering
uplicate matching thresholds	Define safe Customer/Offering/Business match rules and user confirmation thresholds.	Engineering + Product
t dashboard handoff threshold	Confirm whether partially prepared larger tasks may hand off before completion and under what useful-state criteria.	Product
onymous abuse controls	Rate limits, bot protection and task/session caps.	Security + Engineering
capability registry source of truth	Confirm which production registry/configuration is authoritative for TaskDefinitions, Agent tools, engines and capability availability so Goal Resolver cannot rely on LLM knowledge alone.	Product + Engineering
omposite-goal orchestration	Confirm whether MVP stores composite paths only inside TaskSession or introduces a lightweight GoalPlan object. Do not add a new persistent entity unless production discovery shows it is necessary.	Product + Engineering
homepage quick tasks at launch	Recommended: Create Invoice; Create Quotation; Create Business Plan; Prepare Proposal. Confirm entitlement behaviour for Proposal.	Product
homepage goal ordering	Start with product strategy/demand assumptions, then use activation and conversion analytics to reorder without changing the core architecture.	Product + Growth
nding first-value threshold	Confirm the minimum reliable Funding Readiness result that can be shown before dashboard handoff and which inputs are mandatory versus progressive.	Product + Funding Readiness owner

25. Vibe-Coding / Claude Code Implementation Handoff

This PRD should be implemented through the existing EnterprateAI Claude Code governance process. Claude Code must perform production discovery and produce a Reuse / Extend / New mapping before changing schemas or creating services. The implementation must prefer current authentication, business, catalogue, Business Operations, dashboard, event and agent mechanisms over parallel structures.

1. Load this PRD together with Shared Architecture, Adaptive Dashboard, Agentic Orchestration and the owning feature PRDs.

2. Inspect production code and map every conceptual entity/API to actual files, routes, schemas and services.
3. Return conflicts, missing production capabilities and proposed minimal extensions before writing code.
4. Implement one vertical slice first: homepage -> task session -> verify -> resume -> invoice draft -> dashboard review.
5. Run unit/integration/E2E and regression tests; capture changed files and migration implications.
6. Review the diff with a second LLM/ChatGPT against this PRD before commit.
7. Implement goal resolution as registry-driven capability mapping. The LLM may infer the business outcome, but it must not declare an executable capability unless production configuration confirms it.
8. Add explicit automated tests for direct feature requests, composite goals (including “I need more customers”), partial-support requests, true no-match, low-confidence clarification, credit behaviour and capability-hallucination prevention.

Do not let Claude infer architecture from the PRD alone The PRD defines product behaviour. Production code remains the implementation source of truth. New tables, APIs or services require evidence that an equivalent production capability does not already exist.

26. Definition of Done - Version 2

- Homepage presents a central goal-driven chat/task launcher without replacing the rest of the marketing site.
- A visitor can begin a supported task before signup and retain context through verified account creation.
- Invoice first-task journey works end-to-end into a review-ready output on the Adaptive Dashboard.
- Confirmed onboarding information writes to the correct existing business records without duplicate silos.
- Chat, structured forms, engines and agent execution have explicit boundaries and provenance.
- Zero-credit onboarding/basic essential-tool behaviour and 50-credit AI behaviour are enforced server-side.
- Missing information, unsupported capabilities, failures and recovery states are explicit and non-fabricated.
- Existing homepage login, Marketplace Claim, Business Operations and dashboard journeys remain functional.
- Analytics measure goal-to-first-output activation, not signup volume alone.
- Additional tasks can be added through versioned TaskDefinitions without rebuilding the onboarding framework.
- Business-outcome goals that are not literal feature names are resolved through the Goal Resolver and can continue into direct, composite or partial supported paths.
- The LLM cannot advertise an executable capability unless it is validated against the registered production capability/task/tool source of truth.
- The “I need more customers” scenario passes end-to-end: honest boundary, relevant supported path, minimal questions, zero-credit routing and no guaranteed outcome.
- Homepage presents business-owner problems/goals as the primary discovery layer and Quick Tasks as a secondary route, while retaining free-text input.
- “I need funding” routes end-to-end into a truthful Funding Readiness first-value experience with context preserved through verification.
- Business-goal, quick-task and free-text entry modes all converge on the same Goal Resolver, TaskSession, canonical Business and Adaptive Dashboard.
- Every displayed homepage goal is backed by a production-validated support path and no suggestion overpromises an external outcome.
- Analytics can compare goal-led versus task-led acquisition, activation, first value, retention and conversion.

27. Final Product Model

VISITOR ARRIVES

↓

HOMEPAGE ENTRY: POPULAR BUSINESS GOAL / QUICK TASK / FREE TEXT

↓

ONE ENTERPRATEAI ASSISTANT EXPERIENCE

↓

GOAL RESOLUTION: UNDERLYING OUTCOME + REGISTERED CAPABILITY CHECK

↓

DIRECT TASK / COMPOSITE GOAL PATH / PARTIAL-SUPPORT GUIDANCE / TRUE NO-MATCH

↓

SUPPORTED TASK SESSION + MINIMUM CONTEXT + PERSISTENCE CHECKPOINT

↓

VERIFIED USER + AUTHORISED BUSINESS CONTEXT

↓

KNOWN DATA REUSE + MISSING DATA ONLY

↓

CANONICAL RECORDS + DETERMINISTIC ENGINES + REGISTERED TOOLS

↓

AGENT PREPARES SUPPORTED WORK / ASSESSMENT / ACTION PATH

↓

ADAPTIVE DASHBOARD WITH USEFUL OUTPUT

↓

REVIEW / APPROVE / CONTINUE / REASSESS

↓

ONGOING BUSINESS INTELLIGENCE + OPERATIONS

Experience principle EnterprateAI should feel as if the workspace forms itself around the problem the business owner is trying to solve or the job they want completed. The homepage speaks in business outcomes first, then translates those outcomes into supported intelligence, tools and workflows underneath.

28. Source and Architecture Alignment

- EnterprateAI Shared Architecture and Feature Integration Blueprint v0.1 - reuse system of record, separate stage/scope/role/entitlement, deterministic ownership, shared events and modular extension.
- Adaptive Dashboard PRD v1.0 - dashboard remains the concise post-acquisition priority/action surface; Startup, Small Business and Multi-Entity pathways are not recreated by onboarding.
- Agentic Orchestration PRD v1.0 - one Assistant/Agent experience, registered tools, approvals, guardrails, audit and recovery.
- Marketplace Business Index & Claim GTM PRD + UX Flows - separate acquisition route that converges on the same canonical Business/workspace and dashboard architecture.
- Launch Readiness, Funding Readiness and engine specifications - owning services retain their criteria, scoring, evidence and deterministic logic; onboarding only routes and collects required inputs.

- Goal Resolver addition - outcome interpretation is an onboarding concern; capability execution remains owned by existing registered TaskDefinitions, Agent tools, engines, Business Operations modules and entitlement/approval services. This update does not create a parallel execution architecture.