

EnterprateAI Launch Readiness Product Requirements

Version 0.1 | 30 September 2026 | Product and engineering review

1 Purpose and intended outcome

Launch Readiness Check helps a founder or business owner determine what must be completed before launching a business or offering. It turns existing business information, evidence and intelligence outputs into an explainable readiness assessment, critical blockers and a practical action sequence.

The core journey is create a launch initiative, confirm reused data, complete relevant evidence, assess readiness, resolve blockers and reassess. The user makes and records the launch decision. A positive assessment must never automatically publish an offering, take payments or mark the business as launched.

This PRD follows the Shared Architecture and Feature Integration Blueprint v0.1 and shares assessment, evidence, action and forecast contracts with the Funding Readiness PRD v0.1. It extends the existing ontology domains and reuses validated engine services. Production schema and API mapping remains an implementation prerequisite.

Readiness indicates preparation for a defined launch scope and date against a named checklist. It does not guarantee demand, profitability, regulatory compliance or business survival. The system must display commercial risks separately from preparation progress.

2 Users and scope

Pathway	Launch use case	Required experience
Startup	First commercial launch	Guided setup, clear blockers, forecast-based cash and delivery preparation
Small Business	New service or offering	Reuse operating data while isolating incremental launch assumptions and resources
Multi-Entity	Oversight across businesses	Scoped initiative comparison, blocker counts and drill-down without averaging readiness

Proposed MVP supports UK service businesses and software businesses whose delivery and pricing model is supported by a validated calculation profile. It covers an initial business launch and a new service/offering launch within an existing business. New markets, physical products and inventory-intensive models are later profiles. A software subscription model must not inherit one-off service formulas unless their applicability is explicitly validated.

This PRD proposes a modest extension to the blueprint's initial-business-only first release: new service launches use the same initiative structure. Product must confirm whether that second journey ships with MVP or follows immediately afterwards. The data design supports both from the start.

MVP includes initiative setup, data reuse, manual evidence and attachments, a scoped launch checklist, deterministic scoring, blockers, dependencies, actions, an assessment report, dashboard summaries and a recorded launch decision. It uses existing scenario capability where validated. Agent-assisted planning, automatic document extraction, external publishing and market-specific launch templates are later capabilities.

The feature is a readiness service, not a general project management system. It links to Business Plan Generator, Idea Validation, catalogue and Business Operations where those tools already support the required action. Users can complete evidence manually without connecting an accounting system or building a full business plan.

3 User journey and interface

1. Select a business and choose Check Launch Readiness. Create or resume an initiative with a name, launch type, offerings, audience, target date, geography and intended launch volume.
2. Define scope: limited pilot or commercial launch. Record which customers, services, channels and delivery commitments are included. MVP uses commercial-launch criteria; pilot mode is available only when its separate profile has been implemented.
3. Review reused profile, offerings, pricing, demand evidence, forecast and capacity inputs. Display their source and date. Offer corrections as explicit updates or initiative-only assumptions.
4. Work through six short sections: customers and offer; pricing and economics; cash; delivery; operational prerequisites; sales and launch plan. Save progress and allow unknown answers without inventing values.
5. Run the assessment. Present classification, assessment date, evidence coverage, known preparation score, blockers, business risks and three next actions. Show unknown and stale results explicitly.
6. Resolve an action by updating its source record or attaching evidence. Reassess to see the effect; marking an action complete alone cannot change the score.
7. Test a supported scenario, such as launch delay, lower demand or higher delivery costs. Show assumption changes and comparison to an immutable baseline.
8. Record the launch decision: proceed, defer or cancel. When launch actually occurs, an authorised user records the actual date and scope. Retain the as-of assessment and any acknowledged open risks.

Screens: initiative list; setup; guided evidence review; results; actions; scenarios; history and launch decision. Keep section progress separate from readiness. Use plain-language questions and optional explanations; technical identifiers belong in developer interfaces. Support keyboard navigation, mobile layouts and status labels that do not depend on colour.

4 Functional requirements

ID	Requirement	Priority
LR01	Create, edit, archive and resume business-scoped launch initiatives	MVP
LR02	Link offerings and reuse relevant business data with provenance	MVP

ID	Requirement	Priority
LR03	Select a versioned profile and determine checklist applicability	MVP
LR04	Capture evidence, forecasts, prerequisites, milestones and dependencies	MVP
LR05	Assess preparation with explicit blockers, confidence and missing data	MVP
LR06	Generate prioritised, deduplicated actions and retain history	MVP
LR07	Export an assessment report and launch checklist	MVP
LR08	Display entity-scoped dashboard summaries	MVP
LR09	Record launch decisions and actual launch without automatic external execution	MVP
LR10	Compare supported launch scenarios without changing live data	Conditional MVP
LR11	Delegate draft planning or evidence gathering through registered agents	Later

LR10 ships only after scenario inputs and calculations are validated. If deferred, the assessment remains functional and scenario controls state that the capability is unavailable. LLM narrative must not substitute for calculated scenarios.

Initiative lifecycle: draft, preparing, launched, cancelled, archived. Assessment execution and readiness classification are separate states. Deferral updates target date and records a decision event while the initiative remains preparing. Archiving preserves the preceding lifecycle state and assessment history. Reopening a cancelled initiative requires an explicit recorded action. Launching a new version of an offering creates a new initiative linked to the earlier one.

5 Data and ontology mapping

Structure	Ownership	Required fields or links
LaunchInitiative	Decision and Simulation domain	Business ID, type, scope, offerings, geography, target date, owner, status, profile version, revision
LaunchMilestone	Initiative-scoped detail	Outcome, due date, owner, prerequisite references, completion evidence

Structure	Ownership	Required fields or links
LaunchPrerequisite	Initiative-scoped detail	Criterion code, applicability, required-by date, critical flag, evidence and dependency links
Business and Offering	Existing Identity and Commercial domains	Stage, model, country, offer description, pricing and active version
EvidenceItem and EvidenceLink	Shared evidence service	Source, claim, effective date, reviewer, verification status and expiry
Forecast and ForecastLine	Shared Financial domain	Periods, opening cash, receipts, payments, assumptions and baseline/scenario version
Capacity and Dependency	Existing Operational domain	Available units/hours, commitments, resource bottlenecks, backup arrangements
Assessment and results	Shared assessment service	Input snapshot, ruleset/engine versions, criteria, gates, metrics, confidence and recommendations
LaunchDecision	Initiative-scoped audit record	Decision, actor, timestamp, assessment reference, scope, rationale and acknowledged risks

Reuse shared ActionItem and DocumentArtifact records from the architecture blueprint. Do not create separate launch-only evidence, action or forecast databases. Funding and launch may reference the same forecast or evidence version, but each owns its assessment and criteria.

Initiative assumptions do not overwrite canonical business data. Material source edits require the existing write permission and invalidate affected current assessments. Each offering, attachment, action and forecast link must belong to the authorised business. A portfolio view never bypasses those checks.

Store dates with business timezone context and unambiguous calendar semantics. Monetary values require currency and decimal precision. Capacity inputs require units and periods; hours cannot be compared with customers or transactions without an explicit conversion model. Persist immutable input snapshots for reproducibility.

6 Economics cash and capacity rules

Assess the planned launch volume against supported pricing and cost inputs. Reuse contribution margin, fixed-cost and capacity computations from the validated engine profile. Negative contribution must be displayed as a commercial risk. It becomes a blocker when the selected profile requires a positive contribution model and no documented funded pilot exception applies; MVP commercial launch does not use pilot exceptions.

For a new offering, distinguish incremental revenue and costs from shared business costs. Record the allocation method and avoid counting shared costs twice. Shared team capacity is available capacity after existing

commitments, not total headcount multiplied by nominal hours. Where several offerings share a constrained resource, aggregate their demand against that resource before assessing feasibility.

The MVP cash projection covers at least six monthly periods from the assessment month and at least three full months after the planned launch, whichever is longer. This is a proposed product requirement subject to pilot validation. Reuse a longer funding forecast where available. The launch target must fall within the forecast horizon.

Opening cash plus cash receipts and financing receipts minus cash payments gives closing cash. Carry closing cash into the next period. Include setup costs, operating payments and dated commitments. Proposed funding belongs in a scenario until received or separately evidenced as committed, with its expected receipt date. A favourable Funding Readiness result does not establish available funds.

The baseline must identify first negative-cash month, minimum cash and cash available at launch. A confirmed projected cash shortfall within the required horizon is a launch blocker unless a documented, executable mitigation is incorporated into the assessed baseline. A statement that funding will probably arrive is insufficient. Record any committed financing conditions and flag uncertainty.

Capacity feasibility requires planned demand to fit confirmed available delivery capacity in matching units and periods. For constrained launch volume, the user may reduce commitments and reassess. Overload cannot disappear solely because an action to hire someone exists; the input must reflect supported availability by the required date.

Resolve the source-engine boundary conditions, example arithmetic and break-even-months issue identified in the blueprint before reuse. Return typed unavailable results for insufficient inputs. Do not calculate stability from fabricated revenue history for a pre-revenue business.

7 Readiness profile and proposed rubric

The rubric is a pilot design, not a validated predictor of commercial success. The initial profile is commercial service launch. Each criterion has two checks, recorded as passed, failed or unknown with evidence. Both passed gives 100; one passed and one failed gives 50; both failed gives 0. Any unknown check makes the criterion unknown. All eight criteria apply in this profile.

Criterion and weight	Check A	Check B
C1 Customer and demand 15	Target segment, problem and dated demand evidence documented	Evidence relevance and limitations explicitly reviewed
C2 Offer and pricing 15	Launch scope, price, delivery terms and customer promise confirmed	Supported cost and margin inputs validated and implications reviewed
C3 Cash preparation 20	Required cash projection is complete and reconciled	Setup and operating commitments are included with evidence or labelled assumptions
C4 Delivery capability 15	Planned demand fits evidenced available capacity	Critical suppliers, people and systems are available with contingency where needed

Criterion and weight	Check A	Check B
C5 Operational process 10	Enquiry/order to delivery and payment process is documented	A representative end-to-end dry run has recorded results and resolved critical failures
C6 Applicable prerequisites 10	Applicable launch obligations and prerequisites are identified and reviewed	Required permissions or prerequisites are evidenced as completed by launch
C7 Route to market 10	Initial acquisition channel, audience and launch message are defined	Channel execution has an owner, budget and measurable target
C8 Launch control 5	Launch milestones, date and owners are agreed	Success measures, incident response and post-launch review date are recorded

Demand evidence may include research, interviews, enquiries, letters of intent, pilots or sales. Preserve evidence type and strength. A reviewed research item can establish preparation without establishing strong demand; show the Viability findings and commercial uncertainty separately.

Qualitative checks require an authorised user's explicit confirmation and rationale. A founder can self-review in MVP; label this self-attested. Independent review is optional and separately identified. AI can propose extracted facts and interpretations but cannot independently verify a claim or pass a prerequisite.

C6 uses maintained jurisdiction/sector profiles. For MVP, users record and review their applicable prerequisites, source references and completion evidence. The system must not claim its checklist is exhaustive or certify compliance. Where applicability is unresolved, C6 remains unknown. A prerequisite may be not applicable only with a recorded rationale; the criterion itself remains applicable and requires a reviewed applicability determination. No legislation-specific rule ships without verified sources and an assigned maintenance owner.

8 Blockers classification and confidence

Launch blockers operate independently of aggregate score. Proposed blockers are B1 a confirmed cash shortfall in the required horizon; B2 demand exceeding confirmed delivery capacity; B3 a known required launch permission or prerequisite not completed by its required date; B4 a failed critical end-to-end delivery/payment test without resolution; and B5 non-positive contribution under the MVP commercial-service profile.

Every blocker stores trigger inputs, source, severity, explanation and remedy. Unknown inputs needed to evaluate a blocker make its gate unknown, not passed. A missing mandatory prerequisite record is insufficient evidence; a known required permission that has not been obtained is a confirmed blocker. A checklist action being complete cannot override these rules.

When every criterion and gate is known, calculate readiness as the weighted sum of criterion scores; weights total 100. Evaluate unrounded scores and round only the display. With an unknown criterion or gate, return a null overall score and show criterion results and weighted evidence coverage.

Classification precedence: any confirmed blocker gives not_ready; otherwise any unknown mandatory check or gate gives insufficient_evidence; otherwise below 60 gives not_ready; 60 to below 80 gives conditionally_ready;

80 or above gives ready. A ready result is worded “Prepared for the defined launch scope against this checklist.” Unresolved non-blocking gaps and commercial risks remain visible.

Evidence coverage is the total percentage weight of known criteria; gate completeness is also shown separately. Confidence is low if material sources are absent, stale or conflicted; medium if complete/current but relying materially on self-attestation; high only with corroborating records or documented independent review of material factual claims. Forecast assumptions remain forecasts at every confidence level.

Proposed freshness defaults: cash 30 days; launch forecast confirmation 30 days; price/cost and capacity confirmation 30 days; market evidence review 180 days. Required permissions use their stated validity dates and must cover the launch date. Dry-run evidence relates to the assessed process version; a material change invalidates it. Near-launch reminders propose a final check seven days before target date and again on material change. These configurable defaults require pilot validation.

9 Actions dependencies and launch decisions

Prioritise confirmed blockers, then unknown mandatory evidence, then other gaps. Within each group sort by required-by date, prerequisite order, descending criterion weight and stable action ID. Show three next actions with a complete list available. Do not assign arbitrary predicted revenue or score gains to actions.

Milestone dependencies must be acyclic and scoped to the initiative. A successor may be prepared while a prerequisite is incomplete, but cannot be marked satisfied when its completion requires that prerequisite. Missing dates remain visibly unspecified; the system does not invent deadlines. Changing the target date recalculates proposed action dates and asks the user to review changes before replacing manually set dates.

Action states follow the shared model: open, in_progress, awaiting_evidence, completed, dismissed. Repeated assessments update existing unresolved actions for the same criterion/blocker rather than duplicating them. Evidence acceptance and reassessment determine readiness. Dismissing an action records a reason and leaves the criterion unchanged.

The decision screen shows the selected assessment date, freshness, scope and open risks. Proceed, defer and cancel record intent only. Actual launch is a separate user-recorded event. Users may record what occurred even when the assessment was not_ready; preserve that result and record the acknowledged blockers without portraying them as resolved. No UI control represents authorisation to bypass an external legal requirement.

After actual launch, preserve the historical launch assessment and show a post-launch review action. Changing operating stage is an explicit business-profile update, not an automatic consequence of launching one new offering. Historical startup records remain available.

10 Outputs dashboards and connected features

The assessment report contains initiative and business identity, launch scope/date, profile/rules version, classification, score where available, evidence coverage, confidence, blocker details, risks, assumptions, actions and evidence references. Export an editable checklist and PDF readiness report through the existing document service. Each export identifies its assessment snapshot and as-of date.

Startup dashboard prioritises launch blockers, readiness and next action. Small Business dashboard shows the selected new-offering initiative and incremental cash/capacity exposure. Multi-Entity lists authorised initiatives, dates and blockers; no score averaging. Multiple initiatives require a clear selected initiative and a count/link to the others. Default to the nearest upcoming active target date; undated initiatives appear after dated ones.

Summary contract: business_id, initiative_id, launch_type, scope, target_date, lifecycle_status, assessment_id, profile_version, classification, score, evidence_coverage, gate_completeness, confidence, freshness, blocker_count and next_action_id. Null means unavailable. The dashboard must not recalculate scores.

Idea Validation supplies demand and economics evidence where current and applicable. Business Plan Generator supplies referenced assumptions and supporting narrative. Catalogue supplies offerings/pricing; Business Operations supplies process configuration and relevant records. A saved invoice or contract alone is not proof that a complete launch process has been tested. Funding Readiness can supply the funding case reference and shared forecast, but its score never clears a launch cash blocker.

Future agent tools include get_launch_initiative, list_launch_gaps, propose_launch_plan, request_launch_assessment and prepare_launch_checklist. Agents use existing application APIs and shared permission controls. They may draft actions and scenarios but cannot independently verify evidence, clear blockers, change the actual launch date or publish an offering. MVP must work without an agent runtime.

11 API events and operating requirements

Logical operations: create/list/read/update/archive initiative; link offerings and evidence; manage prerequisites/milestones; request/read assessment; manage actions; record decision/actual launch; generate/download report. Use current business-scoped APIs and map names to production conventions during engineering design.

Updates require a revision token; stale writes return a conflict and preserve user edits for resolution. Assessment requests include initiative revision and idempotency key. Persist the input snapshot before engine execution. If live inputs change during a run, keep the historical result and mark it stale. Invalid dates must produce field validation errors rather than dashboard rendering failures.

Events include LaunchInitiativeChanged, EvidenceUpdated, AssessmentCompleted, AssessmentInvalidated, LaunchDecisionRecorded and LaunchRecorded. Include event ID, business scope, schema version and correlation/causation IDs. Publish committed changes reliably and deduplicate consumers. Coalesce updates and prevent reassessment loops. Critical changed evidence invalidates the current result immediately even when recalculation is delayed.

Capabilities: launch.read, launch.edit, launch.assess, launch.review, launch.decide and launch.export. Map them to existing roles. Enforce business and attachment permissions for normal requests, jobs, portfolio queries and downloads. Revoked access applies to queued work. Uploads use existing file validation, malware handling and retention policies. Do not log raw documents, financial values or personal details in analytics.

Proposed engineering budgets: 95th-percentile normal reads below two seconds and deterministic assessment within ten seconds after inputs are available; narrative/report generation remains asynchronous. Validate budgets against declared load before release. LLM failures must not hide valid deterministic results. Jobs expose failure reason, retry/cancel state and recoverable progress; retries must not duplicate artifacts or charges.

12 Acceptance examples

ID	Given and when	Expected result
LT01	All criteria score 100 and every blocker gate passes	Score 100; ready; no automatic launch

ID	Given and when	Expected result
LT02	C3 unknown and no confirmed blocker	Null score; 80 percent criterion coverage; insufficient_evidence
LT03	All criteria known, score 90 and a cash shortfall exists	not_ready; B1 visible despite high preparation score
LT04	Criterion/gate unknown and a required permission is confirmed missing	not_ready; blocker and unknown evidence both shown
LT05	Forecast funding is uncommitted and arrives after a cash gap	Baseline gap remains; scenario cannot clear it
LT06	Existing work uses 80 of 100 available hours; launch needs 30	B2 overload for 10 hours; total nominal capacity cannot conceal it
LT07	Pre-revenue startup has research but no trading history	Research type visible; no invented stability or revenue history
LT08	A pricing or delivery-process version changes after testing	Affected evidence/assessment becomes stale and requires review
LT09	User marks blocker action complete without new evidence	Blocker remains until relevant inputs and reassessment resolve it
LT10	User changes launch date beyond permission validity	Applicability/validity rechecked; expired permission cannot pass
LT11	A scenario changes price, demand or timing	Baseline/live data unchanged; derived results use versioned formulas
LT12	Another business ID or attachment is supplied	No cross-entity access through detail, jobs, exports or aggregates
LT13	Duplicate job/event arrives	One logical execution per idempotency key; no duplicate actions
LT14	User records actual launch with unresolved blockers	Actual event saved with acknowledgement; historic readiness unchanged
LT15	Engine succeeds but narrative generation fails	Deterministic findings and actions remain usable
LT16	A milestone dependency would create a cycle	Save rejected with a clear dependency error

Also test exact score thresholds 60 and 80, unknown gate with full criterion coverage, zero demand, zero capacity, zero/negative contribution, invalid dates, unsupported profiles, mixed units/currencies, simultaneous edits and revoked permissions. Confirm an existing business's operating stage remains unchanged when one offering launches.

13 Delivery measurement and release gates

Work packages: production inventory and profile review; initiative/evidence/dependency persistence; validated calculation and readiness rules; guided UI/dashboard/reporting; pilot and staged rollout. Use additive migrations and feature flags. Preserve old endpoints and data; rollback disables the new flows without deleting drafts or history.

Engineering must resolve applicable source-engine discrepancies, verify shared capacity allocation and document the forecast mapping before scoring release. Product and domain reviewers approve the pilot rubric, freshness policies, blocker wording and first supported business models. Product confirms whether new-service launch and scenarios ship in MVP. The owner of each sector/jurisdiction profile is responsible for maintaining its external requirement sources.

Pilot with startup founders and operating businesses launching services. Independently review sample cases, compare classifications, inspect false-ready cases and identify confusing questions. Refine the versioned rubric before public scoring. Test existing business operations and dashboard journeys for regressions.

Measure case-to-assessment completion, time to first actionable result, reused-input percentage, missing-evidence rate, blocker resolution before target date, reassessment frequency, action usefulness feedback, assessment failures and cost per successful run. Track recorded launch and post-launch review completion as user outcomes without claiming that the feature caused commercial success. Set targets after baseline measurement.

Public release requires representative acceptance fixtures, no unresolved cross-entity access defects, deterministic formula agreement, reproducible snapshots, a recoverable migration/flag plan and reviewed user wording. A formal automated compliance profile additionally requires verified rules and a maintenance owner; a generic self-reviewed checklist must remain labelled accordingly.

14 Traceability and implementation handoff

Architecture references: Shared Architecture and Feature Integration Blueprint v0.1, decisions AD-01 to AD-09 and sections 4, 5, 7 to 13. Reuse the Funding Readiness PRD v0.1 contracts for forecasts, evidence, actions, documents and assessment history, with feature-specific criterion and blocker definitions in this PRD.

Ontology references: Business/BusinessProfile, Offering/PricingModel, DemandEvidence, financial records, CapacityProfile/CapacitySnapshot, Process/Dependency, Assessment/AssessmentMetric/AssessmentScore, RiskFlag/Recommendation and Scenario/Simulation entities in Updated Business Ontology. Engines: Service Idea Viability for supported economics; Survival for cash metrics; Stability for available operational risk evidence; Growth for supporting market context; Real-Time Simulation for isolated scenarios.

Do not require all four engines to return a score before launch assessment can run. Compute applicable metrics and preserve unavailable states. Do not silently renormalise the master intelligence score or introduce a new Fragility Index formula. The Launch Readiness score is its own preparation measure.

Engineering handoff includes field ownership and source mappings, schema migrations, API/event contracts, capability mapping, versioned profile configuration, calculation fixtures, input-validation cases and rollout steps. Design handoff includes all named screens, accessible mobile states, missing-evidence treatment, blocker details, decision recording and portfolio drill-down.

Open decisions to close before affected implementation: first supported model profiles; inclusion of new-service launch; availability of scenario execution; who can record review exceptions; subscription entitlements and

limits; and support ownership for external prerequisite profiles. The next linked PRD is the Adaptive Dashboard, followed by Agentic Orchestration.